

Shaurya Pathak

Santa Clara, CA · shauryapathak24@gmail.com · 562-330-8614 · github.com/shaurya-pathak

Experience

NVIDIA · Santa Clara, CA

Software Development Engineer in Test · Mar 2025 - Present

- Architected automated debugging orchestration across hundreds of release workflows using AI-driven analysis, Celery, RabbitMQ, and Python, reducing false-positive reports by 40%.
- Built multi-platform validation infrastructure for GPU driver stacks, benchmarking latency, throughput, and floating-point precision across varied workloads.
- Engineered CI/CD regression pipelines that reduced validation runtime by 60% while expanding cross-architecture coverage.

NVIDIA · Santa Clara, CA

Data Engineer - NeMo · Aug 2024 - Mar 2025

- Designed ETL pipelines and analytical dashboards supporting LLM training and evaluation across research teams.
- Developed 15+ statistical measures for annotation bias, linguistic diversity, and dataset consistency.
- Built automated validation workflows that made model evaluation more reproducible.

NASA Jet Propulsion Laboratory / City of Los Angeles · Remote

Student Researcher · Sep 2021 - Dec 2024

- Built convolutional LSTM models to forecast PM2.5 from satellite imagery and sensor data, achieving RMSE of 1.6 micrograms per cubic meter without ground stations.
- Engineered an AWS inference pipeline using Glue, Lambda, and S3 to generate hourly pollution forecasts.
- Integrated static and real-time traffic data into simulation frameworks for congestion and urban-planning analysis.

AllCare Therapies · Cerritos, CA

Software Engineer Intern · Jun 2024 - Aug 2024

- Developed a serverless ML classification pipeline with AWS Lambda and Step Functions, automating 90% of document processing.
- Improved data-ingestion and model-inference throughput by 3x through pipeline refinements.

Education

University of California, Riverside · B.S. Computer Science with Business Applications · Dec 2024

GPA 3.82 (major 3.88) · Chancellor's and Dean's List · Undergraduate research assistant

Selected Projects

Get Me to a Human · Operator-supervised outbound voice agent that navigates IVRs, uses DTMF, and grounds support outcomes in the other party's words.

Kaivra · Declarative animation engine that renders structured JSON or YAML to stills, video, and web previews, with a local MCP authoring workflow.

Glyde · Local-first generation and hosting layer for native Slidev decks with reusable interactive components and visual checks.

Technical Skills

Python, SQL, Bash, C++ · Pandas, NumPy, SciPy, scikit-learn, PyTorch, TensorFlow, Celery, RabbitMQ · AWS, Linux, Docker, Git, Ansible, CI/CD, CUDA